

Anchored Conditionally-Gated Control for Wheeled Bipedal Balance and Disturbance Recovery

Van Thuong Nguyen^a

^aIndependent Researcher, Viet Nam

ABSTRACT

Wheeled bipedal robots face a trade-off: the stiffness that buys precise standing shrinks the push-recovery envelope, while the integral action that removes equilibrium bias winds up during disturbances. The Anchored Conditionally-Gated Controller (ACC) is a two-channel torque assembly whose core mechanism is a proximity-gated anchor: a position integral confined to a 15 cm neighbourhood and enabled by an asymmetric envelope follower ($\tau_{\text{attack}} \approx 23$ ms, $\tau_{\text{release}} \approx 1.5$ s) that separates quiet stance from post-push ringdown. It is evaluated on a 10-DOF, 8.1 kg wheeled biped in MuJoCo (500 Hz physics, 100 Hz control) through a factorial ablation, an eight-direction push sweep, and a delay screen resolved at 2 ms. ACC holds 0.294 ± 0.015 mm sagittal repeatability ($N=10$), a 59-fold improvement over a proportional-only baseline, while recovering from omnidirectional pushes from 82 N in the weakest direction to a 107 N median. The ablation separates the mechanisms: gated damping produces the steadiness, the anchor integral produces the accuracy, and the gate itself is inert in quiet stance and decisive under disturbance. A closed-form model predicts the residual 27.0 mm parking offset to within 0.92 % across a fortyfold gain sweep. Six classical baselines and a full-state LQR all fall. Flight recovery from a 100 cm drop and per-leg adaptation to a 10–20 cm curb are demonstrated at $N=10$. Conditional gating, not higher gain, buys precision and disturbance rejection together; actuator latency is the sharpest deployment constraint, the push envelope collapsing by 45 % between 6 ms and 8 ms.

ARTICLE HISTORY

Received ***

Accepted ***

KEYWORDS

Wheeled biped robot;
conditionally-gated control;
anchored standing;
proximity-gated integral;
push recovery; flight
recovery; terrain
adaptation.

1. Introduction

Wheeled biped robots need millimeter-level standing precision, omnidirectional push recovery, and terrain adaptation—but no single classical controller integrates all three: LQR handles balance (Klemm et al., 2019), WBC posture (Klemm et al., 2020), MPC terrain (Wang et al., 2025), and RL locomotion (Hwangbo et al., 2019). The difficulty is fundamental. Lowering k_p widens the capture basin at the cost of balance stiffness; omitting integral action leaves a 1.3 Nm equilibrium bias that sustains a persistent, predominantly *sagittal* limit cycle regardless of k_p ($N=10$: 69.7 mm peak-to-peak sagittally against 13.4 mm laterally), on the one axis the wheels can drive; yet a global integral winds up during pushes.

We present the Anchored Conditionally-Gated Controller (ACC), which addresses both dilemmas through *conditional activation*. The core mechanism is a **proximity-gated anchor**: a position integral confined

to a 15 cm neighborhood, enabled by an **asymmetric envelope follower** ($\tau_{\text{attack}} \approx 23$ ms, $\tau_{\text{release}} \approx 1.5$ s) that distinguishes quiet stance from post-push ringdown. Coupled with a **gated damping boost**, ACC achieves 0.294 ± 0.015 mm repeatability on the sagittal axis under idealized conditions (clean sensors, 0 ms delay, $N=10$), at push performance equal to or better than a proportional-only baseline carrying no integral action at all. Absolute accuracy on that axis is a separate quantity: the robot parks 27.0 mm from the commanded home and then holds that standoff to 0.009 mm.

A factorial ablation under a single protocol (Table 1) separates three effects the architecture couples: quiet-stance *steadiness* comes from the gated *damping* terms rather than from the anchor integral, which buys *accuracy* instead; the proximity gate and the asymmetric envelope are provably inert during quiet stance, so only a displacement experiment can measure them; and the

pitch safety gate is load-bearing with no disturbance applied. The gates are enabling conditions that make an otherwise inadmissible damping and integral configuration safe, not the proximate source of the precision. The paper is positioned within *structured classical priors* rather than learned control: the object of study is a hand-designed, interpretable torque assembly whose downstream role is to supply the nominal base action of a residual-learning stack (Johannink et al., 2019; Silver et al., 2018). Classical controllers are therefore the comparison of record: six of them—four 4-state TWIP variants and two 6-state coupled variants—all fail, and a full-state LQR linearized from the plant itself survives two orders of magnitude longer and still falls, so platform complexity, not model order or servo lag, dominates the failure. A preliminary model-free PPO (Schulman et al., 2017) run is retained as a *difficulty probe* and bounds nothing about what RL attains here.

The main contributions are:

- (1) A **proximity-gated anchor** with asymmetric envelope follower achieving sub-millimeter idle precision without sacrificing omnidirectional push survival.
- (2) A **two-channel conditionally-gated torque assembly** with flight recovery and per-leg terrain adaptation as independent extensions, each ablated at $N=10$.
- (3) A **factorial ablation** (additive L0→L3 + subtractive S0→S5) against six classical baselines, plus two control experiments that test the comparison itself: a full-state LQR derived from the plant rather than a reduced model, and a sweep of the coupled model’s one empirical constant.

2. Related Work

2.1. Wheeled Bipedal Platforms and Classical Control

Ascento (Klemm et al., 2019) demonstrated two-wheeled jumping with LQR balancing, later extended to WBC with kinematic loops (Klemm et al., 2020); SKATER (Wang et al., 2025) and comparable wheeled-leg platforms pair balance control with separate yaw, height, roll or terrain loops. Whleaper (Zhu et al., 2024) is closest to our morphology (10 DOF, three hip DOFs per leg); it applies LQR to balanced sliding and a separately trained PPO policy to walking, *switching*

between the two by locomotion mode rather than composing them on a shared control output. Composition of a classical prior with a learned residual on one output is the residual-RL construction (Johannink et al., 2019; Silver et al., 2018), applied recently to wheeled-biped balance (Zhang et al., 2025); ACC supplies the classical prior in that construction rather than replacing the learned part. A recent review (Liu et al., 2024) notes the absence of a unified classical solution.

Coupled pitch-roll LQR is standard on rigid two-wheeled platforms, where torque acts directly on the wheels and there are neither articulated legs nor leg-height dynamics. Our 6-state LQR reproduces that structure on a legged morphology and fails; the 4-state Direct Torque LQR, which removes the servo path entirely, also fails, so actuation lag alone cannot account for the outcome (§5.4). These results are empirical, not a proof that LQR is structurally insufficient.

Cascade control and whole-body control via hierarchical inverse dynamics or task-priority quadratic programs (Herzog et al., 2016) use fixed-gain nested or strictly prioritized loops without conditional gates, and gain scheduling varies parameters continuously with one scheduling variable rather than responding to qualitatively distinct physical regimes; the task-priority QP we benchmark standalone and as an assist layer on ACC (§5.5) fails to resolve the precision–recovery trade-off for exactly that reason. Split-range, selector and override architectures (Åström & Hägglund, 2006) already blend continuously in process control; ACC’s contribution is the synthesis of physically-motivated gating surfaces—spatial proximity, temporal velocity envelope, pitch safety interlock—derived from wheeled-biped failure modes rather than from generic process-variable limits.

2.2. Disturbance Rejection, Anti-Windup, and ACC’s Distinction

DOB (Ohnishi et al., 1996) and ADRC (Han, 2009) provide temporal disturbance estimation but cannot distinguish a push from post-push ringdown—both occupy the same frequency band. Classical anti-windup—back-calculation, conditional integration and the unified observer-based framework (Åström & Rundqvist, 1989)—triggers on *actuator saturation*; ACC’s anchor instead uses **spatial, physics-conditioned gating**, confining integration to a 15 cm neighborhood and dis-

tinguishing quiet stance from ringdown by a velocity envelope, which replaces the estimation-speed-versus-noise trade-off with a time-domain attack-versus-release trade-off. Push recovery via capture point (Englsberger et al., 2015) and flight-phase control (Roscia et al., 2023) is complementary, and terrain adaptation is usually exteroceptive or learned (Lee et al., 2020); ACC extends proprioceptive contact-point estimation to wheeled bipeds. Hyon & Cheng (2006) made a humanoid passive to unmeasured external forces by gravity compensation with optimal contact-force distribution; ACC meets the same requirement by switching regime on proprioceptive gating variables rather than redistributing contact forces.

3. Robot Platform

3.1. Robot Morphology and Simulation

The robot (Fig. 1) is a ten-DOF wheeled biped with five actuated joints per leg: hip roll (HR), hip yaw (HY), hip pitch (HP), knee (KN), and drive wheel (WH). Mass: 8.1 kg; thigh: 0.26 m; shin: 0.28 m; wheel radius: 0.06 m; hip width: 0.23 m. Torque limits: 30 Nm (HR, HY, WH), 150 Nm (HP, KN). MuJoCo (Todorov et al., 2012) simulation runs at 500 Hz ($\Delta t_{\text{sim}} = 0.002$ s); ACC and every classical baseline run at 100 Hz ($\Delta t_{\text{ctrl}} = 0.01$ s, 5 physics substeps), with the baselines additionally measured at their 50 Hz design rate and the PPO reference reported at the 50 Hz rate it was trained at. All experiments use a 400 Nm/s torque rate limit and raw torque commands direct to actuators. Integrator: implicitfast; cone: pyramidal (elliptic for wheels, $\mu_s=1.2$, $\mu_k=0.8$ body); armature 0.008 kg·m² (wheels) and 0.02 kg·m² (leg joints); contact solref={0.002,1} (body), {0.005,0.7} (wheels). Softening solref [0] to 0.02 (tire-on-asphalt) increases planar idle RMS by ~60 % to ~1.2 mm, so every sub-millimeter idle figure is conditional on stiff ground contact.

Height convention

CoM-z, the whole-body centre-of-mass height above ground, is what ACC commands and regulates; the nominal standing pose is $h_{\text{CoM}}=0.404$ m. **Base-z**, the torso root-body height (qpos [2]), sits 13.2 cm above it in that pose, varying by less than 1 cm over the postures used here. All ACC results are in CoM-z; the PPO reference (§5.4) and the WBC baselines (§5.5) command base-z and are labelled accordingly, since substituting one convention for the other injects the full

13.2 cm offset as a reference error. ACC's posture map is calibrated at 0.354 m and 0.454 m and interpolated between them; outside that band joint targets are linearly extrapolated and bounded to [0.254, 0.554] m, and it is that ± 10 cm bound the per-leg height split (§4.6) operates against.




Figure 1. Robot joint layout with world frame (right-handed, Z up). Both wheel axes lie along world X and the wheel links are separated by 0.271 m along it, so the rolling—and therefore sagittal—direction is world Y and world X is the lateral (track) axis; all axis-resolved results follow that convention (Sec. 5.1). Each leg: hip roll (HR), hip yaw (HY), hip pitch (HP), knee (KN), wheel (WH).

4. ACC Formulation

4.1. Two-Channel Torque Assembly

ACC assembles a PD balance core, a proximity-gated anchor integral and posture regulation into two torque channels—wheel torque ($\tau_w \in \mathbb{R}^2$) and leg-joint torque ($\tau_q \in \mathbb{R}^8$)—with flight and terrain as independent extensions (Fig. 2):

$$\begin{aligned}\tau_w &= \tau_{\text{balance}} + g_{\text{anchor}} \tau_{\text{anchor}} + g_{\text{flight}} \tau_{\text{flight}}, \\ \tau_q &= \tau_{\text{posture}}(h^{\text{cmd}}) + g_{\text{terrain}} \Delta \tau_{\text{posture}}(h^{\text{cmd},L}, h^{\text{cmd},R}).\end{aligned}\quad (1)$$

Every gate g_* except the flight gate is a *smoothstep* function of a physical condition (proximity, quietness, attitude, terrain disparity) rather than a discrete switch:

$$\text{ss}(x; a, b) = 3t^2 - 2t^3, \quad t = \text{clip}\left(\frac{x-a}{b-a}, 0, 1\right). \quad (2)$$

Equation (2) is the standard Hermite blend, C^1 in x with zero derivative at both knees. Continuity matters concretely: a hard threshold injects a torque step equal to the full gated component, which at 100 Hz under a 400 Nm/s rate limit reaches the plant as an unmodelled impulse. The one exception is g_{flight} , a discrete load-based hysteresis with asymmetric engage/release timers, because contact loss is a binary physical event.

Wheel torques control sagittal balance and yaw; leg-joint torques control posture and per-leg height adaptation. The two channels are coupled only through the plant—no term in τ_w reads a leg-posture state and none in τ_q reads the sagittal balance state—which is what allows flight recovery and terrain adaptation to be added as independent extensions without retuning the anchor. Unlike the industrial override architectures of §2.2, which select *among parallel controllers* on process-variable limits, ACC gates *components within a single torque assembly* on physical quantities of three kinds—spatial, temporal and safety-interlock—each chosen because a specific measured failure mode of this platform is expressible in it.




Figure 2. ACC two-channel torque assembly. Physical conditions gate each component via continuous smoothstep functions (dashed).

4.2. Sagittal Balance Core (τ_{balance})

The sagittal balance component is always active and acts entirely in the wheel-torque channel:

$$\tau_{\text{balance}} = \begin{bmatrix} \tau_{\text{com}} + \tau_{\text{diff}} \\ \tau_{\text{com}} - \tau_{\text{diff}} \end{bmatrix}, \quad \begin{aligned} \tau_{\text{com}} &= \tau_{\text{pitch}} + \tau_{\text{damping}} + \tau_{\text{pos}}, \\ \tau_{\text{diff}} &= \tau_{\text{lat}} + \tau_{\text{yaw}}, \end{aligned} \quad (3)$$

with common-mode and differential components

$$\tau_{\text{pitch}} = k_p \cdot \theta + k_d \cdot \dot{\theta}, \quad (4)$$

$$\tau_{\text{damping}} = -k_{\text{vel}} \cdot v_{\text{sag}}, \quad (5)$$

$$\tau_{\text{pos}} = -k_{\text{pos}} \cdot \text{clip}(\Delta x, \pm 0.10 \text{ m}), \quad (6)$$

$$\tau_{\text{lat}} = k_{p,\text{roll}} \cdot \phi + k_{d,\text{roll}} \cdot \dot{\phi}, \quad (7)$$

$$\tau_{\text{yaw}} = k_{p,\text{yaw}} \cdot e_{\psi} + k_{d,\text{yaw}} \cdot \omega_z, \quad (8)$$

where θ is torso pitch, v_{sag} the sagittal body velocity, Δx the sagittal position error from the latched home position, ϕ the roll angle and e_{ψ} the yaw error.

The P position term (k_{pos} , clipped to ± 0.10 m, saturating at ± 4 Nm) keeps the robot within $\pm 4 - 6$ cm of home but leaves a 1.3 Nm equilibrium bias sustaining a perpetual limit cycle, which is what motivates the anchor integral. Saturation-triggered anti-windup does not address it, because the critical failure is *state-dependent*: during a push, position error grows to 10–30 cm while the wheels remain within their 20 rad/s range, so a saturation trigger never fires and ACC uses spatial confinement instead. Pitch stiffness is fixed at $k_p=50$; a scheduled variant softening it to 35 during pushes was tested and rejected (S0 vs. S1, Table 2).

4.3. Proximity-Gated Anchor ($g_{\text{anchor}} \tau_{\text{anchor}}$)

The anchor mechanism is the core contribution. It adds two terms to the wheel-torque channel that are active only when the robot is near home *and* genuinely at rest:

$$\begin{aligned} \tau_{\text{anchor}} &= K_i \int g_{\text{prox}} \cdot g_{\text{env}} \cdot (-\Delta x) dt \\ &\quad + g_{\text{boost}} \cdot K_{\text{boost}} \cdot (-v_{\text{sag}}), \end{aligned} \quad (9)$$

where the gates are continuous smoothstep functions equal to 1 when the robot is “at home and quiet” and transitioning smoothly to 0 otherwise:

$$g_{\text{prox}}(\Delta x) = 1 - \text{ss}(|\Delta x|; 0.05 \text{ m}, 0.15 \text{ m}), \quad (10)$$

$$g_{\text{env}}(v_{\text{sag}}) = 1 - \text{ss}(\text{EMA}(|v_{\text{sag}}|); 0.25 \text{ m/s}, 0.50 \text{ m/s}), \quad (11)$$

$$\begin{aligned} g_{\theta}(\theta, \dot{\theta}) &= [1 - \text{ss}(|\theta|; 2^\circ, 12^\circ)] \\ &\quad \times [1 - \text{ss}(|\dot{\theta}|; 2^\circ/\text{s}, 15^\circ/\text{s})], \end{aligned} \quad (12)$$

$$g_{\text{boost}} = g_{\text{prox}} \cdot g_{\text{env}} \cdot g_{\theta}. \quad (13)$$

Each gate is dimensionless, monotone and C^1 , and three physically distinct conditions must hold simultaneously for the anchor to integrate: spatially near home (g_{prox}), temporally quiet (g_{env}) and attitudinally settled (g_{θ}).

The pitch gate closes when either pitch or pitch rate leaves its band, preventing the damping boost from fighting a gentle sustained push that does not trigger the velocity envelope.

Asymmetric envelope follower. Reliable quiet-stance detection rests on an asymmetric exponentially-weighted moving average (EMA) of $|v_{\text{sag}}|$:

$$\text{EMA}_t = \begin{cases} \alpha_a |v_t| + (1-\alpha_a)\text{EMA}_{t-1}, & |v_t| > \text{EMA}_{t-1}, \\ \alpha_r |v_t| + (1-\alpha_r)\text{EMA}_{t-1}, & \text{else,} \end{cases} \quad (14)$$

with time constants $\tau_{\text{attack}} \approx 23$ ms and $\tau_{\text{release}} \approx 1.5$ s as the primary specification, the discrete coefficients following as $\alpha_a = 1 - e^{-\Delta t / \tau_a} \approx 0.35$ and $\alpha_r \approx 0.007$ at 100 Hz. The asymmetry is essential: a symmetric EMA is *bistable*—post-push oscillation holds it in a middle band where the gate is half-open and the boost cannot collapse the cycle, as the symmetric-EMA ablation of §5 confirms—whereas the fast-attack/slow-release envelope rises to 1 only when oscillation truly decays and drops to 0 within ~ 25 ms of a disturbance. The mechanism is a peak-hold envelope detector; ACC's contribution is coupling it to a spatial proximity gate for integral conditioning. Since $\text{EMA}_0 = 0$ biases $g_{\text{env}} \approx 1$ for ~ 7.5 s at startup, all measurements discard a 3–5 s settle interval.

4.4. Posture and Yaw Stability (τ_{posture})

Posture regulation uses Jacobian-based PD gains scheduled across 23 heights:

$$\tau_{\text{posture}} = \tau_{\text{PD}}^{\text{jac}}(q, \dot{q}, h^{\text{cmd}}) + \tau_{\text{div}}(q_{\text{hy}}^L, q_{\text{hy}}^R) + g_{\text{settled}} \tau_{\text{homing}}(q - q_{\text{ref}}), \quad (15)$$

where τ_{div} damps hip-yaw divergence and τ_{homing} restores nominal posture when upright and settled. Yaw return uses wheel-differential torque.

4.5. Contact-Loss Recovery ($g_{\text{flight}} \cdot \tau_{\text{flight}}$)

When both wheels lose ground contact, ACC detects this from per-wheel normal forces and substitutes a flight attitude controller using the wheels as reaction flywheels:

$$g_{\text{flight}} = \begin{cases} 1, & \text{if } F_z^L, F_z^R < 0.5mg \text{ for } \geq 2 \text{ steps,} \\ 0, & \text{if } \min(F_z^L, F_z^R) \geq 0.5mg \text{ for } \geq 5 \text{ steps,} \end{cases} \quad (16)$$

$$\tau_{\text{flight}} = \text{clip}(2.0\theta + 0.5\dot{\theta} - 0.02\omega_{\text{wheel}}, \pm 2 \text{ Nm}). \quad (17)$$

Forward wheel spin pitches the torso nose-up during descent, and the 5-step release hysteresis prevents balance-torque re-engagement during bounce. The wheel-inverted-pendulum loop this override replaces is itself a reaction-flywheel controller once the wheels are unloaded, so the override's measured contribution is to bound landing-attitude variance near the upper end of the drop envelope rather than to enable recovery (§5.3).

4.6. Per-Leg Ground Adaptation (g_{terrain})

When the two wheels rest at different ground heights—a curb straddle, an oblique ledge—a single torso height command forces the torso to roll by the step angle. ACC instead maintains an *independent ground estimate per leg* and issues two height commands. The estimate is taken from the wheel contact-point height, which is exact at any body tilt for a rigid wheel, and is updated only while that wheel carries real load:

$$\begin{aligned} \hat{h}_{\text{gnd}}^i &\leftarrow \hat{h}_{\text{gnd}}^i + \alpha_g (z_{\text{contact}}^i - \hat{h}_{\text{gnd}}^i), \\ F_z^i &\geq 0.2mg, \quad i \in \{L, R\}, \end{aligned} \quad (18)$$

with $\alpha_g = 0.35$ per control step to reject contact chatter. The $0.2mg$ load test rather than mere geometric touching is deliberate: a lightly touching wheel in the act of unloading injected a 2.9 cm phantom ground step on the 15 cm curb. An unloaded wheel freezes its estimate, the one exception being the unambiguous signal that its ground is gone—the wheel centre descending *below* its own believed ground:

$$\delta^i > 3 \text{ mm} \Rightarrow \dot{\hat{h}}_{\text{gnd}}^i = -(v_0 + k_s \delta^i), \quad (19)$$

where $\delta^i = \hat{h}_{\text{gnd}}^i - z_{\text{contact}}^i$, $v_0 = 0.5$ m/s and $k_s = 50 \text{ s}^{-1}$. The proportional term matters: a fixed slew converged in 0.30 s, by which time a 20 cm dismount had already tipped the robot past recovery, whereas (19) converges within two control steps, so the leg extends *during* the fall. When both wheels are airborne both estimates track the mean wheel height.

The step $\Delta h = \hat{h}_{\text{gnd}}^L - \hat{h}_{\text{gnd}}^R$ is then absorbed by the legs rather than the torso. The split is kinematic, not a linear height offset: with $L(\cdot)$ the forward-kinematic leg length and L^{-1} its inverse over the calibrated posture range, the leg on the higher ground is shortened by the full step,

$$L_{\text{high}} = L(h^{\text{cmd}}) - |\Delta h|, \quad (20)$$

$$h^{\text{high}} = L^{-1}(L_{\text{high}}), \quad h^{\text{low}} = h^{\text{cmd}}, \quad (21)$$

both clamped to the calibrated envelope extended by 10 cm at each end. To first order this is the familiar

$h^{\text{cmd}} \mp \Delta h/2$ split, but the $h \mapsto L$ map is markedly non-linear near full squat, the linear proxy erring by 12 cm at the lower extrapolation limit. If L_{high} falls below the shortest reachable leg, ACC raises the *whole* robot until the short leg fits and reports the remainder as an expected torso roll $\arctan(-\varepsilon/w_{\text{track}})$ (ε the unspanned step, $w_{\text{track}} = 0.278$ m), which keeps a legitimate straddle from being read as an incipient fall. Because the $h \mapsto L$ table is itself extrapolated at the extremes, a purely feedforward split left a systematic 6° torso roll on the 10 cm curb; a closed-loop levelling integrator trims it by the *measured* roll at 0.15 m/(rad s) with ± 6 cm of authority, integrating only when both wheels are loaded, the step is genuine ($|\Delta h| \geq 2$ cm) and the estimate is settling (< 0.15 m/s); the first condition keeps a 180° turn from winding in an 8.6° bias, the second keeps the transient of a curb climb out of the levelling error.

4.7. Complete Numerical Parameter Set

The core gains are $k_p=50$ Nm/rad, $k_d=3$ Nm/(rad/s), $k_{\text{pos}}=40$ Nm/m (clipped at ± 4 Nm), $k_{\text{vel}}=15$ Nm/(m/s) at damping scale 1.5, $k_{\text{roll}}=55$ Nm/rad, $k_{\text{yaw}}=1.0$ Nm/rad; the anchor uses $K_i=4.0$ Nm/(m s) capped at 2 Nm with a 5×10^{-3} per-step leak, boost scale 5.0, g_{prox} band 0.05–0.15 m, g_{env} band 0.25–0.50 m/s, g_{stab} bands $2-12^\circ$ and $2-15^\circ \text{ s}^{-1}$; flight uses P/D = 2.0/0.5 with 0.02 Nm/(rad/s) wheel damping and 2/5-step engage/release at $F_z=0.5mg$ (0.2mg for terrain). The full set (50+ gains) is in the supplementary material.

5. Experimental Validation

ACC and all six classical baselines are evaluated at 500 Hz physics / 100 Hz control under a 400 Nm/s per-joint torque rate limit, so the headline comparison is rate-matched by construction; the PPO reference is reported at the 50 Hz rate its policy was trained at. Each classical baseline was additionally measured at its own 50 Hz design rate: the fall rate is 100 % in all twelve cells and five of the six survive *longer* at 50 Hz, so reporting them at 100 Hz is the rate-matched and not the flattering choice. Seeds are deterministic: 20260727+trial_id for idle standing, 20260728 $\times$ 100+rep for the push ablation, and {0, 42, 123} for the classical baselines.

5.1. Anchored Standing

5.1.1. Metrics and axis convention

“Standing precision” names two quantities that are not interchangeable. With $\mathbf{p}(t) \in \mathbb{R}^2$ the horizontal base

position over a window of length T , $\mathbf{p}_0$ the commanded home and $\bar{\mathbf{p}}$ the trial mean,

$$A = \left(\frac{1}{T} \int_0^T \|\mathbf{p}(t) - \mathbf{p}_0\|^2 dt \right)^{1/2}, \quad (22)$$

$$R = \left(\frac{1}{T} \int_0^T \|\mathbf{p}(t) - \bar{\mathbf{p}}\|^2 dt \right)^{1/2} \quad (23)$$

define the *accuracy* A , which retains any DC parking offset, and the *repeatability* R , which removes it. The two differ by more than an order of magnitude on this platform, so both are reported for every configuration. Each is resolved along the sagittal axis, the lateral axis and the full plane, because the two horizontal axes are actuated by different means: the **sagittal** axis is world y , the rolling direction and the only one on which wheel torque produces ground force, while the **lateral** axis is world x , held passively by the wheel contacts and the roll restoring moment of the track width. A dynamic check confirms the assignment (scripts/verify_body_axes.py): a 90 N impulse along y is absorbed at a peak wheel speed of 52 rad/s leaving 4.8 mm of residual offset, whereas the same impulse along x leaves a permanent 23.2 mm offset. The anchor integral and the gated damping boost both live in the wheel-torque channel, so both act on the sagittal axis only.

5.1.2. Protocol

Every row of Table 1 was re-measured under one protocol: each trial runs 25 s from the nominal standing pose with randomised initial conditions ($\sigma=0.005$ rad per joint, clipped to limits; $\sigma=1$ mm on root height), the first 5 s discarded as settling (long enough to cover the $\text{EMA}_0=0$ transient to within 1 %) and A , R computed over the remaining 20 s relative to each trial’s own $t=0$ pose. A trial is a fall if $|\theta| > 46^\circ$ or root height < 0.15 m; confidence intervals are Student- t . The protocol measures the ability to *stay* put; coming back is measured in Sec. 5.2.

5.1.3. Result

ACC holds **sub-millimetre sagittal repeatability**, $R_{\text{sag}} = 0.294 \pm 0.015$ mm (95% CI [0.283, 0.305] mm) across ten randomised trials with no falls—the figure of merit for the platform, since sagittal is the actuated inverted-pendulum axis. Precision is assessed from the relative spread *between* configurations under an identical protocol rather than against a simulator noise floor: removing a single mechanism (S2) moves R_{sag} to 4.772 mm with non-overlapping SDs. Against the P-only baseline L0, which sustains a 69.7 mm peak-to-peak sagittal limit cycle, the like-for-like

Table 1. Idle standing precision across the full ablation ladder. All rows re-measured under the single protocol of Sec. 5.1; values are mean $\pm$ SD over $N=10$ trials.

Configuration		R_{sag} (mm) sagittal rep.	R_{xy} (mm) planar rep.	R_{lat} (mm) lateral rep.	$\bar{s}$ (mm) sag. offset	Survival
Additive ladder (bottom-up)						
L0	P-only (PD, no position, no integral)	17.461 ± 0.258	17.845 ± 0.279	3.675 ± 0.246	-44.8	10/10
L1	+ position homing (± 0.10 m)	20.397 ± 0.017	20.414 ± 0.019	0.834 ± 0.066	-30.4	10/10
L2	+ ungated integral K_i	25.026 ± 0.039	25.056 ± 0.043	1.212 ± 0.110	-27.5	10/10
L2.5	+ proximity gate only (no envelope, no boost)	25.198 ± 0.017	25.227 ± 0.020	1.205 ± 0.099	-27.6	10/10
Proposed controller						
L3/S1	ACC (full anchor, fixed $k_p=50$)	0.294 ± 0.015	0.788 ± 0.069	0.730 ± 0.073	-27.0	10/10
Subtractive ablation (top-down from ACC)						
S2	– gated damping boost	4.772 ± 0.069	4.824 ± 0.067	0.703 ± 0.069	-26.8	10/10
S3	– asymmetric envelope (symmetric EMA)	0.294 ± 0.015	0.788 ± 0.069	0.730 ± 0.073	-27.0	10/10
S4	– proximity gate (integral always on)	0.294 ± 0.015	0.788 ± 0.069	0.730 ± 0.073	-27.0	10/10
S5	– pitch safety gate g_θ	—	—	—	—	0/10
Single-mechanism isolation (ACC with one term disabled)						
M1	$K_i=0$ (anchor integral removed, gates intact)	0.089 ± 0.035	0.733 ± 0.072	0.727 ± 0.071	-31.2	10/10
M2	damping terms reverted to their L2 values	25.198 ± 0.017	25.227 ± 0.020	1.205 ± 0.099	-27.6	10/10
M3	$K_i=0$ and boost = 0	3.973 ± 0.130	4.037 ± 0.128	0.710 ± 0.069	-31.1	10/10

100Hz control, 20s window after a 5s settle, clean sensors, 0ms delay. R = repeatability (Eq. (23)); $\bar{s}$ = mean signed sagittal parking offset; accuracy on an axis is that offset and the repeatability in quadrature, so the A columns are omitted (the identity holds to 0.003 mm on every surviving row). Sagittal is world y , lateral world x ; the anchor integral and the damping boost act sagittally only, so R_{sag} and $\bar{s}$ discriminate between configurations while R_{lat} measures a passively held axis. Produced by `scripts/collect_idle_ladder.py`; raw output `outputs/paper_verification/idle_ladder.json`.

improvements are $17.461 \rightarrow 0.294$ mm sagittal (a factor of **59.4**), $17.845 \rightarrow 0.788$ mm planar (**22.6**) and $3.675 \rightarrow 0.730$ mm lateral (**5.0**)—the largest gain on the hard axis and the smallest on the easy one, the correct signature for a wheel-torque controller, since every term ACC adds beyond L0 produces ground force along the rolling direction and none can produce lateral force at all.

Accuracy and repeatability are distinct on this axis. ACC parks at $\bar{s} = -27.0$ mm from commanded home and holds that standoff to 0.009 mm SD across trials, so planar accuracy is $A_{xy} = 26.994$ mm, a $1.8\times$ improvement over L0 against the $59\times$ gain in repeatability. Its magnitude is fully accounted for: instrumenting the sagittal torque balance predicts the offset in closed form from a *leaky* anchor integral whose ~ 2 s forgetting horizon fixes its DC gain at the finite $k_{I,\text{dc}}=7.96$ Nm/m—so a constant load leaves a finite error by construction—and a 1.44° error in the feedforward pitch reference, giving 24.88 mm against 24.99 mm measured and holding across a $40\times$ gain sweep and a five-point height sweep to within 0.92 %. The leak also buys the steadiness: a $10\times$ smaller λ cuts the offset to -16.31 mm but degrades window jitter twelvefold. Table 1 publishes a deliberately chosen point on that trade-off.

5.1.4. Which mechanism produces the precision

The lower blocks of Table 1 isolate the contributions, and three of the four findings contradict the attribution that the gate structure invites.

(i) *Idle steadiness comes from the damping ladder; the integral buys accuracy instead.* M1 disables the anchor integral ($K_i=0$) with every gate in place, and its *repeatability* is slightly better than ACC's (0.089 mm) while its *accuracy* is worse by 4.2 mm; conversely M2 keeps the integral and every gate but reverts the two velocity-damping coefficients to their L2 values, reproducing the L2/L2.5 limit cycle exactly. Steadiness is therefore attributable to the damping terms in two stages: raising the velocity-damping scale takes R_{sag} from 25.2 mm to 3.97 mm (M3), and enabling the gated boost takes it to 0.294 mm, a combined factor of 86.

(ii) *The damping boost acts sagittally, as it must.* S2 degrades sagittal repeatability sixteenfold while leaving lateral repeatability statistically unchanged (0.703 mm vs. 0.730 mm, overlapping CIs), the only outcome consistent with a boost inside the wheel-torque channel.

(iii) *The proximity gate and the asymmetric envelope are inert during quiet stance.* Ablating either (S4, S3) is *bit-identical* to ACC on every idle metric—the base never leaves the 5 cm inner band and the velocity envelope never rises into the attack branch—so both are correctly evaluated only in Table 2.

(iv) *The pitch safety gate is*

required even to stand. S5 falls in 10/10 trials during quiet stance with no push applied: removing g_θ leaves the damping boost free to act on the pitch state at large excursions, and the resulting positive feedback diverges from the ± 0.005 rad initial perturbation alone. The anchor's *gates* are therefore what make an aggressive damping and integral configuration admissible without windup or divergence, while the damping terms they admit are what deliver the steadiness.

5.2. Omnidirectional Push Recovery

Push recovery is evaluated by polar sweep: 10–160 N impulses of 7 control steps applied at the torso in 8 world-frame directions ($\theta = \pm 90^\circ$ sagittal, $\theta = 0^\circ/180^\circ$ lateral), each threshold located by binary search to 5 N, with survival defined as torso pitch $< 46^\circ$ for 17 s post-push. $F_{\min}$ is the maximum magnitude survived in every direction and F_{med} the median.

The envelope is anisotropic by construction. The two lateral bearings average 98.2 N against 131.3 N for the two sagittal ones, and the weakest bearing of all is the oblique 45° (82.3 N, 23 % below F_{med}). The asymmetry follows from the morphology rather than from tuning: the wheels can drive the contact patch back under the CoM sagittally but produce no lateral ground force, so raising the lateral bearing requires a centralised lateral task (§5.5), not a gain change.

The envelope is also chiral, and the chirality is in the controller. Superposed on the anisotropy is a left–right asymmetry: 135° survives 40.3 N more than -135° (Fig. 3). The plant cannot produce it—mirroring the model about the sagittal plane reproduces its mass, inertia and geometry to 0.05 mm and 1×10^{-7} —and a harmonic fit localises the violation in the mirror-forbidden term ($b_2 = -16.3$ N against $b_1 = 0.3$ N). The source is the sign convention of the two differential channels: under a mirrored state the commanded torque maps to its own mirror image exactly on hip-pitch, knee and wheel and fails only on hip-roll and hip-yaw, both assembled antisymmetrically—even rather than odd under the mirror map. Mirroring those channels flips b_2 to 15.6 N at unchanged magnitude, while a sham mirror leaves it at -15.7 N. We report the effect without adopting the correction, since it redistributes the envelope without enlarging it ($F_{\min} 82.5 \rightarrow 79.7$ N). Later protocols apply a single impulse along world $+x$, a lateral and below-median bearing (95.6 N), so those experiments are conservative probes.

Table 2. Factorial ablation of ACC components: push-recovery and ringdown metrics. Clean sensors, all rows $N=10$, mean $\pm$ std.

Configuration		$F_{\min}$ (N)	F_{med} (N)	Ringdown
Additive (bottom-up)				
L0	P-only (PD, no pos., no I)	70.0 $\pm$ 1.0	122.1 $\pm$ 2.1	never
L1	+Pos. homing (± 0.10 m)	74.8 $\pm$ 0.8	130.5 $\pm$ 1.2	never
L2	+Integral K_i (no gate)	74.8 $\pm$ 1.0	131.1 $\pm$ 1.1	—
L2.5	+Prox. gate only (no EMA, no boost)	75.3 $\pm$ 0.8	131.9 $\pm$ 2.4	—
L3	+Full anchor (gate, EMA, boost)	82.3$\pm$2.2	106.7 $\pm$ 7.0	9.0 s
Subtractive (top-down from ACC)				
S0	ACC + scheduled k_p (rejected)	81.0 $\pm$ 2.2	109.2 $\pm$ 7.6	9.0 s
S1	ACC (canonical, fixed $k_p=50$)	82.3$\pm$2.2	106.7 $\pm$ 7.0	9.0 s
S2	–Damping boost	80.8 $\pm$ 0.9	116.2 $\pm$ 5.3	never
S3	–Asym. EMA (sym. EMA)	81.6 $\pm$ 1.8	106.2 $\pm$ 6.7	never
S4	–Prox. gate (global I)	80.5 $\pm$ 2.7	109.2 $\pm$ 7.6	windup
S5	–Pitch gate g_θ	<10	<10	never

$F_{\min}$ and F_{med} are the minimum and median survivable push over 8 bearings, each located by binary search to 5 N tolerance; idle metrics for the same configurations are in Table 1. The S5 entries are left-censored at the search floor: in all 80 trials the bisection never located a survivable magnitude in [10 N, 160 N]. Welch's unpaired two-sided t -test ($N=10$ per group, Bonferroni-adjusted $\alpha=0.007$): L0 $\rightarrow$ L1 $p < 1 \times 10^{-9}$, $d=5.4$; L1 $\rightarrow$ L2 $p=0.88$; L2 $\rightarrow$ L2.5 $p=0.19$; L2.5 $\rightarrow$ L3 $p < 2 \times 10^{-6}$, $d=4.1$. No subtractive row differs significantly from ACC except S5 ($p < 1 \times 10^{-14}$). At $N=10$ the 95% CI half-width is $\approx 0.72 \times \text{std}$, so sub-2 N differences are not resolvable. Each trial starts from a freshly constructed controller context. Produced by `scripts/replicate_ablation_n10.py`; raw output `outputs/paper_statistics/`.




Figure 3. Omnidirectional push envelope of ACC—the same run as row S1 of Table 2. Solid curve: per-bearing mean survivable push over $N=10$ repetitions; shaded band ± 1 SD; radial axis linear in force from 0 N. The envelope spans 82.3–144.8 N and is both anisotropic and chiral.

Table 2 reports the factorial ablation. Additively, L1 (homing) adds 4.8 N over L0 ($p < 1 \times 10^{-9}$), L2 (global integral) adds nothing over L1, and L3 (full anchor) adds 7.5 N over L2 ($p < 1 \times 10^{-6}$); the proximity gate contributes nothing to push *margin* on its own (L2→L2.5, n.s.), consistent with its role being windup prevention rather than capture-basin enlargement, while the full anchor is simultaneously the only configuration in the ladder that settles. Subtractively the striking result is how flat the ablation is: every single-mechanism removal except the pitch gate leaves $F_{\min}$ within 1.8 N of ACC and none is significant, so capture basin is set by the anchor as a whole rather than by any one of its four gates—including gain scheduling ($k_p=50$ fixed costs nothing measurable, S0) and the damping boost, which dominates idle precision yet leaves push margin indistinguishable from ACC's (S2). The one mechanism the push data *do* isolate is the pitch gate: removing it (S5) leaves the robot unable to survive the 10 N search floor in any of 80 trials. Post-push ringdown decays within 9 s under ACC, whereas P-only oscillates indefinitely because the 1.3 Nm equilibrium bias sustains the limit cycle; the symmetric-EMA ablation exhibits the bistable failure the asymmetry is designed to prevent, and the ungated integral winds up during the push and prevents settling entirely.

5.3. Flight Recovery and Terrain Adaptation

Both extensions run at the same rate limit and control rate as the core with no re-tuning of the anchor; all cells are $N=10$ except the 150 cm drop ($N=5$).

Contact-loss recovery. ACC recovers to a settled stand from a 100 cm free drop (10/10, $24.0 \pm 1.2^\circ$ peak pitch after touchdown) and from a 50 cm ledge drive-off (10/10, $27.1 \pm 2.0^\circ$), with the 60 cm drop ($18.2 \pm 1.0^\circ$) and 20 cm ledge ($19.6 \pm 0.3^\circ$) correspondingly milder; throughout, survival is the *strict settle* verdict (flat pitch/roll window, tail height error ≤ 12 mm, residual speed ≤ 0.05 m/s), not merely staying upright. The two flight ablations are more nuanced than the design rationale suggests: removing the flight *attitude* PD is a real but modest loss (peak pitch 27.8° on the 100 cm drop, one fall in ten), while removing the airborne wheel-torque override entirely does not degrade either condition and in fact *lowers* peak pitch to $13.0 \pm 0.5^\circ$ at 10/10, because free wheels driven by the pitch-feedback loop are themselves an effective reaction flywheel. The override's benefit appears only at the

upper end of the envelope: at 150 cm both configurations settle 2/5, but the override holds peak pitch to $34.0 \pm 2.1^\circ$ against $53.6 \pm 35.0^\circ$, so it **bounds landing-attitude variance near the envelope limit** rather than enabling the recovery.

Per-leg terrain adaptation. Here the ablation is unambiguous. With per-leg adaptation, straddle roll on the one-wheel curb is $2.5 \pm 0.04^\circ$ (10 cm), $3.5 \pm 3.2^\circ$ (15 cm) and $14.9 \pm 0.1^\circ$ (20 cm), all 10/10. Replacing the split with a symmetric posture—the pre-adapter behaviour—falls at *every* tested curb height, 0/10 in all three cases, the torso rolling past 60° within 0.8 s of curb entry. Disabling the closed-loop levelling trim alone leaves the robot upright on the 10 cm curb but degrades the 15 cm curb to $11.1 \pm 1.6^\circ$ roll with 6/10 settled. At 20 cm the limit is joint authority, not split geometry: (21) commands 0.350 m and 0.150 m leg lengths, spanning the full step with *zero* kinematic residual, while the compressed leg's knee reaches 2.58 rad against a 2.70 rad limit—the measured 14.9° roll is therefore entirely *tracking* error at the shortest gravitational moment arm the PID gains were tuned for.

Oblique ledge descent: a straddle regime, not a height threshold. Driving off a ledge whose edge crosses the path at an angle ϕ releases the two wheels a stagger interval $w_{\text{track}} \tan \phi / v$ apart, and only within that interval can the leg split absorb the step. Instrumenting a 30° , 20 cm exit at 0.7 m/s shows the window closing in 0.18 s with 99 mm of calibrated range still unused, so the failure is a race rather than a saturation and is consequently *not* monotone in step height. A 192-trial calibration resolves the boundary as a success-probability field: at 30° the shallow window runs to 15 cm at 100%, closes completely over 22–27 cm, and a deep regime—both wheels released together, the flight path of §4.5 taking over—reopens from 30 cm, peaking at 75%. Widening the straddle regime requires knee range rather than gains.

5.4. Comparison with Classical Baselines and a Preliminary RL Reference

To establish that ACC is not merely a tuned linear controller we compare six classical baselines under identical environment setups: four derived from the 4-state TWIP model (Grasser et al., 2002) (LQR with decoupled roll/yaw PD through a PID servo layer of ~ 0.25 s lag; the same plus a pitch-error integral with conditional anti-windup; a direct-torque variant

re-optimised for the zero-lag plant; and PI+AW, adding a conditional forward-position integral) and two 6-state coupled variants adding the roll states $(\phi, \dot{\phi})$, one through the servo and one commanding torque directly. All weights follow Bryson’s rule (Bryson & Ho, 1975). A model-free PPO run (5.4M steps, single seed, no domain randomisation or hyperparameter search) is reported alongside them as a difficulty probe rather than a tuned RL baseline (Limitation 4).

Table 3. Classical baselines and the preliminary PPO reference vs. ACC, clean sensors. All rows at 100 Hz control except the PPO reference, which is reported at the 50 Hz rate its policy was trained at.

Configuration	Surv. (s)	θ_{RMS} (°)	ϕ_{RMS} (°)	H_{RMS} (mm)	τ_{RMS} (Nm)
LQR, TWIP ($N=20$)	0.32	0.006	24.2	205.8	21.99
LQR + AW ($N=20$)	0.32	0.006	24.2	205.8	21.99
LQR, DT (re-opt.) ($N=20$)	0.60	0.033	27.5	119.2	8.76
6-St. LQR ($N=20$)	0.32	0.006	24.2	205.8	21.99
PI + AW, DT ($N=20$)	0.52	0.181	26.1	171.3	10.42
6-St. LQR, DT ($N=20$)	0.16	0.001	21.2	128.5	14.44
PPO ref. (50 Hz) ($N=20$)	3.20	4.82	3.15	54.2	18.40
ACC ($N=10$)	20.00	0.060	0.010	0.18	3.28

LQR / LQR+AW / 6-St. use the PID servo; the DT rows command torque directly. The three servo rows are identical because that servo’s response dominates before the gain differential between them can act. Survival time is deterministic on all six classical rows (± 0.00 s at $N=20$, seeds $\{0, 42, 123\}$). Raw output `outputs/rate_matched_baselines/results.json`.

All six classical baselines fall within 0.60 s. Removing the PID servo path cuts torque by $2.5\times$ and nearly doubles survival (0.32 s $\rightarrow$ 0.60 s) but *worsens* roll ($\phi_{\text{RMS}}=27.5^\circ$ vs. 24.2°): the recovered authority is spent in the sagittal plane the model describes, not in the lateral one that ends the episode. A 46-configuration retuning sweep leaves the outcome unchanged, as does sweeping the coupled model’s one empirical constant, the hip-roll-to-roll-acceleration coupling $b_{r,h}$, over an $80\times$ range: the servo variant falls at 0.320 s for every value, with roll RMS spread 0.06 %, even though the re-derived roll gain varies $32\times$. That is invariance to the parameter rather than insensitivity of the metric, and its mechanism is saturation—the roll channel is commanded against the ± 0.7 rad hip-roll limit throughout the episode at every value tested. The preliminary PPO reference sustains single-height standing longer (3.20 s) but exhibits an 85 % fall rate across multi-height command transitions; read as a difficulty probe, that is the observation the structured prior is meant to address.

5.4.1. A full-state LQR designed on the plant itself

All six baselines are designed on a reduced model, so their failure is open to the objection that the model, not the method, is what fails. The control experiment that settles this is a regulator derived from the plant directly: a damped Gauss–Newton solve for a true static equilibrium at the commanded height, central differences on the whole held-input control step of the complete 16-DOF, 10-actuator MuJoCo model for the transition Jacobians, and a discrete-time Riccati solve on the 30 states left after removing the two cyclic wheel angles, with Bryson’s rule fixing $\mathbf{Q}$, $\mathbf{R}$ and a uniform control weight r swept over six decades. The design is sound on every diagnostic that bears on it—residual unbalanced force 0.428 N against 79.46 N of body weight, $\mathbf{A}$ of full rank 30 and PBH-stabilisable, Riccati relative residual $\sim 10^{-14}$, closed-loop max $|\lambda|=0.9911$ —and $\kappa(\mathbf{A})\approx 10^{10}$ is a 50 Hz artifact (2.6×10^7 at 100 Hz) from strongly contracting contact modes that obstructs neither the solve nor the feedback.

What the sweep shows is a slow divergence a short horizon hides. Over 20 s the design looks successful ($r=1000$ at 0.65 m: 0 % falls, 11.0 mm height RMSE); at 60 s every configuration falls in 100 % of episodes at 22.2–57.0 s, with planar drift growing monotonically to 0.43–0.63 m—that drift, not attitude, ends them. Nor does one weight cover the band: $r=100$ stands 20 s at 0.60 m and 0.69 m but falls at 0.65 m, $r=1000$ does the reverse, and under random height commands the fall rate runs 45–80 %; the maximum recoverable push is 10.0–20.2 N against ACC’s $F_{\min}=82$ N. Full-state optimal linear feedback about a single equilibrium therefore buys two orders of magnitude of survival over the reduced-order designs and still does not close the problem.

5.5. Whole-Body Control Baselines

Whether a tuned whole-body controller could match ACC was tested over 225 scenarios spanning fixed-height balance, height transitions and push recovery at 20–120 N under cloned initial conditions (481 000 control steps, no gain retuning): a task-priority quadratic-program WBC—minimising CoM height error, torso orientation error, posture damping and wheel regularisation subject to full-body dynamics and contact constraints—used as the *primary* balance controller (W1), and the same WBC blended as an *assist* on ACC via $\tau_{\text{assist}} = \tau_{\text{ACC}} + \alpha(\tau_{\text{wbc}} - \tau_{\text{ACC}})$ (W2). So that the

comparison is not a verdict on our own implementation, the WBC was verified beforehand against executable references: its analytic contact Jacobian against MuJoCo's `mj_jac`, its task stack against four independently selectable weightings, and its height reference against the CoM- z convention of §3.1.

W1. The standalone QP-WBC falls in **225 of 225** scenarios, where ACC falls in none; it holds the robot upright for 0.570 s and spends 97.3 % of every episode fallen, and the failure is *lateral* (roll RMS 93.45° against ACC's 3.71°). Neither implementation quality nor task weighting decides it: a degraded build solving only 85.85 % of steps falls at 0.52 s against the verified build's 0.570 s at 99.99 % solved, and four task-weight modes move the fallen fraction by 0.1 percentage points. The mechanism is that the posture task is a velocity damper, not a position tracker—its reference is rebuilt from the measured joint vector every control step, so it collapses to $-k_{d,\text{posture}}\dot{\mathbf{q}}$ and can slow a postural drift but not reverse one.

W2. The two assist configurations together constitute a gate ablation. With the authority gate *disabled* ($\alpha=0.25$ fixed) the assist arm never falls—ACC carries it—but degrades ACC on every metric but one: pitch RMS from 0.52° to 1.295°, yaw drift by 1.9×, planar drift by 2.0×. With the gate *enabled*, the admitted authority on the one variant where it opens at all is $\bar{\alpha}=0.0061$, 2.5 % of the permitted maximum, and there the gated assist agrees with ACC alone to four significant figures; on the remaining four variants α is identically zero. Roll is the single exception: with the gate disabled it improves systematically from 3.71° to 3.20°, because ACC regulates the lateral axis with decentralised hip-roll PD whereas the WBC uses a centralised orientation task exploiting full mass-matrix coupling. That 18 % of roll costs 2.5× the pitch error and 2.0× the drift, so a centralised lateral task is a credible direction for ACC's weakest axis—the 98.2 N mean lateral bearing of Sec. 5.2—but not through torque-level blending with a posture stack that cannot hold a posture. Compute is not the constraint: the QP solve costs 0.169 ms mean and 0.30 ms p99 (OSQP (Stellato et al., 2020)).

5.6. Architectural Analysis

Robustness to sensor noise and actuator delay. A factorial sweep over four noise levels and three delay levels ($N=20$ per cell, 240 trials) bounds ACC on both axes at once (Table 4). One control period of delay

is benign for stability and expensive for performance: at 10 ms there are zero falls on clean and low noise, but $F_{\max}$ bifurcates from 95.7 N to 50.5 N and idle RMS rises 5.6×. Three control periods is a different regime: at 30 ms the platform falls in 20 of 20 trials on *clean* sensing. Noise alone is survivable up to the medium level; what ACC does not tolerate is the *conjunction*, medium noise with 10 ms of delay costing 2 falls in 20, the only sub-catastrophic falls anywhere in the grid. High noise is unrecoverable at every delay including zero, because the asymmetric envelope cannot distinguish noise from motion—the same mechanism that produces the sharpest single sensitivity in the design, idle repeatability degrading 71-fold from 0.43 mm clean to 30.82 mm under medium noise at zero delay as the envelope gate releases and re-engages the integral. ACC is therefore specified against a filtered velocity estimate, a requirement on the state estimator rather than a limit on the controller.

Table 4. Robustness to sensor noise and actuator delay, $N=20$ per cell. Falls: trials terminated by a fall, out of N . Idle RMS: lateral (world x) CoM std. over 20 s after a 3 s settle, *surviving trials only*. $F_{\max}$: binary-search lateral (world $+x$) push, ± 5 N tolerance, search floor 10 N. Delay is quantised to the control period here; it is resolved five times more finely in the text.

Condition	Falls	Idle RMS (mm)		$F_{\max}$ (N)	
		Mean	$\pm$ Std	Mean	$\pm$ Std
Clean, 0 ms delay	0/20	0.43	0.06	95.7	0.8
Clean, 10 ms delay	0/20	2.43	0.21	50.5	7.9
Clean, 30 ms delay	20/20	— (all fell)		—	
Low noise, 0 ms delay	0/20	5.86	2.53	96.0	1.0
Low noise, 10 ms delay	0/20	10.66	3.70	51.3	8.9
Low noise, 30 ms delay	20/20	— (all fell)		—	
Med noise, 0 ms delay	0/20	30.82	12.13	93.5	5.2
Med noise, 10 ms delay	2/20	35.38	21.63	51.3	13.4
Med noise, 30 ms delay	20/20	— (all fell)		—	
High noise, all three delays	20/20	— (all fell)		—	

Noise ($1-\sigma$): low = 0.01 rad/s angular velocity, 0.01 m/s² gravity, 0.001 rad joint position; med = 5× low; high = 10× low. Delay is a transport lag on the actuator command. High noise fails identically at all three delays, so those rows are merged. Because the idle column averages survivors only, the medium-noise 10 ms row is mildly optimistic and its falls column is the primary metric. Produced by `scripts/collect_robustness_sweep.py`; raw output `outputs/paper_statistics/robustness_sweep.json`.

Resolving the delay axis. The grid quantises delay to the control period, so the knee was located separately with the transport delay applied per *physics* substep (2 ms at 500 Hz), the clean-sensing harness being deterministic. The knee lies between 6 ms and 8 ms: $F_{\max}$ is flat to

within 2.4 % out to 6 ms (93.2–95.5 N), then collapses 45 % to 51.0 N and declines monotonically to the 10 N search floor by 20 ms—a genuine loss of margin rather than a narrow phase-alignment notch, and it sits *below* the 9.5 ms end-to-end budget estimated in §6. Losing the push margin and losing the platform are different thresholds: on quiet stance the standing threshold is 14–16 ms, station held through 14 ms at 10.84 mm idle lateral RMS against 0.43 mm undelayed.

The sweep also answers whether that knee is architectural. Holding the delay at 8 ms and sweeping one profile constant at a time gives not a gently sloped optimum but a sharp *local dip*: the shipped value of each of four constants tested scores 51.0 N while nearly every neighbouring value recovers 78–93 N, which is not how an exhausted margin behaves. Carrying the largest single-knob change forward—sagittal velocity damping doubled to 45 Nm/(m/s), denoted ACC-DH—leaves the controller identical at 0–2 ms and moves the knee outward to 8–10 ms (51.0 → 93.2 N at 8 ms) for +1.4 % idle lateral RMS. We report ACC-DH as an evaluation variant and do *not* adopt it, since the constant it changes also governs driving acceleration and settling.

Gate sensitivity and rate-limit invariance. Finite-difference gradients reveal a stark hierarchy: the release coefficient α_r dominates (sensitivity 220.4, 20× the next and 130× the pitch gates), while the pitch thresholds are effectively insensitive safety interlocks. The three gate failure modes all map onto measured ablation rows of Table 2: stuck-at-1 causes windup and a fall (S4), stuck-at-0 yields the P-only limit cycle (L0), and stuck-half-open triggers the parametric resonance of S5. The 400 Nm/s slew limit is not what sets the push ceiling: over an 8× span lateral $F_{\max}$ is 92.5±1.1 N at 100 Nm/s and 94.4±0.0 N at 200, 400 and 800 Nm/s, while during a 90 N push the commanded slew saturates exactly at each bound—so the ceiling is architectural rather than actuator-limited.

Compound disturbance. Because the posture map is height-scheduled, a commanded transition and a push compete for the same leg-torque budget, so the standard lateral impulse is applied at the midpoint of a 1 s commanded height ramp of size Δh about the 0.404 m nominal, with $\Delta h=0$ as the matched static control ($N=10$ per row). Inside the calibrated band a simultaneous height command is nearly free (a +5 cm rise costs 5.1 % of push margin; a –5 cm squat is *better*

than static, +6.5 %), and the degradation outside it is one-sided: squatting –10 cm costs nothing measurable ($p=0.99$) while rising +10 cm costs 68.7 %, because rising drives the knee toward extension—where the leg Jacobian is worst conditioned—while raising the CoM. Re-parameterising the *same* posture family so achieved height matches the command recovers 78 % ($p=1.5 \times 10^{-4}$) but closes only 35 % of the deficit. Sequential disturbance is absorbed cleanly: a 90 N lateral impulse followed 2 s later, before the first ringdown has decayed, by a 60 N counter-directed impulse is survived **10/10** at episode peak pitch $11.89 \pm 0.51^\circ$, almost all of it from the first impulse—what $\tau_{\text{release}} \approx 1.5$ s is designed to produce, the damping boost still partly engaged when the second impulse arrives.

6. Discussion and Conclusion

Every gate in ACC was added after tracing a measured failure mode of this platform to its root cause: a bistable symmetric envelope, global integral windup, parametric pumping of a scheduled stiffness, hard-leash oscillation and self-referential stiffness. All thresholds were frozen before the ablation sweep, so Table 2 scores a fixed specification rather than a tuned one, and because the gating variables are physical quantities the controller runs without training and is hypothesised—pending hardware validation—to carry a reduced sim-to-real gap relative to a learned policy (Grimming et al., 2020).

Statistical power. Welch’s t -test at Bonferroni-adjusted $\alpha=0.007$ (7 comparisons) makes L0→L1 and L2.5→L3 significant ($p<0.001$; Cohen’s $d=5.4$ and 4.1). Flight and terrain results are $N=10$ with Clopper–Pearson 95% CI [0.69, 1.00]; push ablations resolve ~2 N, and no claim here rests on a smaller difference.

Limitations. (1) *Simulation only.* All results are MuJoCo; hardware validation remains future work. (2) *Deployment timing is the binding constraint.* ACC as tuned requires a 100 Hz loop and an actuator latency clearing the 6–8 ms push-margin cliff of §5.6. Rate independence is not claimed: at 50 Hz with coefficients recomputed for the longer step the platform does not stand, and whether that is excitation of plant modes or phase lag from discretising the asymmetric envelope at $\Delta t=20$ ms is unresolved. (3) *Two calibrated envelopes are non-uniform.* Oblique ledge descent has two survivable regimes separated by a dead band, and upward height commands issued during a push are restricted, a +10 cm rise costing 69 % of $F_{\max}$ while descent is

free at every amplitude tested. (4) *The learning comparison is not resourced, and one classical family is untested.* The model-free PPO run (5.4M steps, one seed, no domain randomisation (Tobin et al., 2017) or hyperparameter search) is a difficulty probe bounding nothing about RL here; the comparison this controller is built for is residual learning *over* the prior. A tuned convex MPC over the linearisation of §5.4.1 is the natural next comparator, since a terminal cost addresses the slow drift divergence that limits the full-state LQR. (5) *Generality across morphologies is undemonstrated.* Every result is for one 10-DOF platform with one set of gains.

Deployment risks. ACC uses raw torque, a 400 Nm/s rate limit and standard sensing (IMU plus encoders). The binding figure is the 6–8 ms latency cliff in push margin (§5.6); the stability threshold is roughly twice as far out (station held through 14 ms), the ordering the architecture predicts, since the gates depend on physical state rather than timing precision. Commission against the 6 ms cliff: an estimated non-RTOS worst-case budget (readout ~1 ms, estimation ~2 ms, computation ~3 ms, actuation ~3.5 ms) totals ~9.5 ms and sits past it, making real-time scheduling a requirement rather than an optimisation. Projected hardware degradation of the 0.294 mm repeatability is 1–3 mm under encoder quantisation, backlash, stiction, tire compliance and IMU noise (cf. Table 4).

Conclusion. We presented ACC, a two-channel torque assembly in which a PD balance core is augmented by a proximity-gated anchor integral, achieving 0.294 ± 0.015 mm idle repeatability on the sagittal axis—a $59\times$ improvement over the proportional-only baseline under one unified protocol ($N=10$, Table 1)—together with omnidirectional push survival from $F_{\min}=82$ N to a $F_{\text{med}}=107$ N median. Absolute accuracy is reported separately and fully accounted for: the robot parks 27.0 mm from commanded home and holds that standoff to 0.009 mm, a closed-form model predicting the offset as the static price of a deliberately bounded integrator memory to within 0.45 %. Six classical baselines and a full-state LQR derived from the plant itself all fall, and adding integral action to LQR-optimal feedback does not prevent it. ACC tolerates moderate noise, is insensitive to the slew bound over an $8\times$ span, and handles contact loss and terrain adaptation without training; its sharpest deployment constraint is actuator latency.

The contribution is a *structured classical prior*: an interpretable controller whose gating encodes the plant’s measured failure modes, intended to supply the nominal base action of a residual-learning stack. Future work is hardware validation, a resourced learning comparison, and extending conditional activation to locomotion.

Code and Data Availability

The controller, the MuJoCo model, every evaluation script referenced above and the raw result files backing each table are available at <https://github.com/thuong180702/Wheeled-bipedal-robot-simulation> (branch main); each results table names its producing script and raw result file in the note beneath it. The extended version of this paper, in the same repository, carries the full derivations, the whole-body-control implementation details and the closed-form parking-offset model.

7. Notes on contributors



Van Thuong Nguyen is an independent researcher based in Viet Nam, working on legged and wheeled mobile robots. His research interests include classical and learning-based control for wheeled bipedal plat-

forms, whole-body control, and structured model-based controllers used as action priors for residual reinforcement learning.

References

- Åström, K. J., & Hägglund, T. (2006). *Advanced PID Control*. Research Triangle Park, NC: ISA—The Instrumentation, Systems and Automation Society.
- Åström, K. J., & Rundqwist, L. (1989). Integrator windup and how to avoid it. In *1989 American Control Conference* (pp. 1693–1698). IEEE.
- Bryson, A. E., & Ho, Y.-C. (1975). *Applied Optimal Control: Optimization, Estimation and Control* (Rev. ed.). Washington, DC: Hemisphere Publishing.
- Engelsberger, J., Ott, C., & Albu-Schäffer, A. (2015). Three-dimensional bipedal walking control based on divergent component of motion. *IEEE Transactions on Robotics*, 31(2), 355–368.

- Grasser, F., D'Arrigo, A., Colombi, S., & Rufer, A. C. (2002). JOE: A mobile, inverted pendulum. *IEEE Transactions on Industrial Electronics*, 49(1), 107–114.
- Grimminger, F., Meduri, A., Khadiv, M., Viereck, J., Wüthrich, M., Naveau, M., ... Righetti, L. (2020). An open torque-controlled modular robot architecture for legged locomotion research. *IEEE Robotics and Automation Letters*, 5(2), 3650–3657.
- Han, J. (2009). From PID to active disturbance rejection control. *IEEE Transactions on Industrial Electronics*, 56(3), 900–906.
- Herzog, A., Rotella, N., Mason, S., Grimminger, F., Schaal, S., & Righetti, L. (2016). Momentum control with hierarchical inverse dynamics on a torque-controlled humanoid. *Autonomous Robots*, 40(3), 473–491.
- Hwangbo, J., Lee, J., Dosovitskiy, A., Bellicoso, D., Tsounis, V., Koltun, V., & Hutter, M. (2019). Learning agile and dynamic motor skills for legged robots. *Science Robotics*, 4(26), eaau5872.
- Hyon, S.-H., & Cheng, G. (2006). Gravity compensation and full-body balancing for humanoid robots. In *2006 6th IEEE-RAS International Conference on Humanoid Robots* (pp. 214–221). IEEE.
- Johannink, T., Bahl, S., Nair, A., Luo, J., Kumar, A., Loskyll, M., ... Levine, S. (2019). Residual reinforcement learning for robot control. In *2019 IEEE International Conference on Robotics and Automation (ICRA)* (pp. 6023–6029). IEEE.
- Klemm, V., Morra, A., Gulich, L., Mannhart, D., Rohr, D., Kamel, M., ... Siegwart, R. (2020). LQR-assisted whole-body control of a wheeled bipedal robot with kinematic loops. *IEEE Robotics and Automation Letters*, 5(2), 3745–3752.
- Klemm, V., Morra, A., Salzmann, C., Tschopp, F., Bodie, K., Gulich, L., ... Siegwart, R. (2019). Ascento: A two-wheeled jumping robot. In *2019 IEEE International Conference on Robotics and Automation (ICRA)* (pp. 7515–7521). IEEE.
- Lee, J., Hwangbo, J., Wellhausen, L., Koltun, V., & Hutter, M. (2020). Learning quadrupedal locomotion over challenging terrain. *Science Robotics*, 5(47), eabc5986.
- Liu, X., Sun, Y., Wen, S., Cao, K., Qi, Q., Zhang, X., ... Ji, A. (2024). Development of wheel-legged biped robots: A review. *Journal of Bionic Engineering*, 21(2), 607–634.
- Ohnishi, K., Shibata, M., & Murakami, T. (1996). Motion control for advanced mechatronics. *IEEE/ASME Transactions on Mechatronics*, 1(1), 56–67.
- Roscia, F., Cumerlotti, A., Del Prete, A., Semini, C., & Focchi, M. (2023). Orientation control system: Enhancing aerial maneuvers for quadruped robots. *Sensors*, 23(3), 1234.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.
- Silver, T., Allen, K., Tenenbaum, J., & Kaelbling, L. (2018). Residual policy learning. *arXiv preprint arXiv:1812.06298*.
- Stellato, B., Banjac, G., Goulart, P., Bemporad, A., & Boyd, S. (2020). OSQP: An operator splitting solver for quadratic programs. *Mathematical Programming Computation*, 12(4), 637–672.
- Tobin, J., Fong, R., Ray, A., Schneider, J., Zaremba, W., & Abbeel, P. (2017). Domain randomization for transferring deep neural networks from simulation to the real world. In *2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)* (pp. 23–30). IEEE.
- Todorov, E., Erez, T., & Tassa, Y. (2012). MuJoCo: A physics engine for model-based control. In *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)* (pp. 5026–5033). IEEE.
- Wang, Y., Chen, T., Rong, X., Zhang, G., Li, Y., & Xin, Y. (2025). Design and control of SKATER: A wheeled-bipedal robot with high-speed turning robustness and terrain adaptability. *IEEE/ASME Transactions on Mechatronics*, 30(2), 1310–1321.
- Zhang, J., Lu, B., Cao, H., & Fang, Y. (2025). Hybrid balance control for wheeled bipedal robot via residual reinforcement learning optimization. In *2025 10th Asia-Pacific Conference on Intelligent Robot Systems (ACIRS)* (pp. 59–66). IEEE.
- Zhu, Y., He, S., Qi, Z., Yong, Z., Qin, Y., & Chen, J. (2024). Wheaper: A 10-DOF flexible bipedal wheeled robot. In *2024 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)* (pp. 11272–11277). IEEE.